

CS 2

Introduction to Programming Methods



Last Time

Networking and Graph Algorithms

- Dijkstra and all...

Final Words on Networking

Notion of ports

- different channels
 - IP header contains add + port

Notion of sockets

- chat over ports
 - listen/bind,recv/send,...

Notion of layers

Application	Protocol	Port
DNS lookup	UDP	53
FTP	TCP	21
HTTP	TCP	80
POP3	TCP	110
Telnet	TCP	23

The diagram illustrates the network stack layers and the flow of data between processes. At the top, a 'Server Process' box contains a sequence of operations: `select()`, `bind()`, `listen()`, `accept()`, `getpeername()`, `recv()`, `process request`, and `send()`. To its right, a 'Client Process' box contains `select()`, `connect()`, `send()`, and `recv()`. Below these, a 'User Application' box is shown. The network stack layers are represented by dashed lines: 7. Application (containing User Application), 6. Transport (containing TCP and UDP), 5. Network (containing IP), and 4. Data Link (Physical). Arrows show data flow from the User Application through the Transport, Network, and Data Link layers to the Physical layer, which is connected to the Network. A small building icon is at the bottom left of the slide.

CS2 - INTRODUCTION TO PROGRAMMING METHODS

25



Today's Topic

Concurrency

- designing and understanding systems that use parallel processing
 - virtually *all* systems in use today use parallel processing, including most safety-critical systems used in cars, airplanes, and medical devices
- touches on fundamental issues
 - predictable *performance*, functional *correctness*
 - of multicore processors, supercomputers, etc.
- just an intro
 - see G. Holzmann's class(es) and CS24 for more




What is Concurrency?

Dealt with sequential algorithms only

- can point to code & say “the program is here”
 - single “thread” of execution; Program Counter

A bit simplistic...

- many programs happen concurrently
 - Word, Photoshop, email, music, ... (called processes)
- even within a program, multiple “threads”
 - Word: interpreting keystrokes, formatting line and page breaks, spell checking...



Real vs. Simulated Parallelism

Computations happen at same physical time

- on different computers
- on multiple processors of a same computer
- on multiple cores of a single processor

Computations *seem* to happen at same time

- when **OS**/**program** switches rapidly between different **programs**/**threads**

Fine and dandy if no shared resources...

- concurrency issues otherwise



Switching between Processes

Processes managed by *kernel*


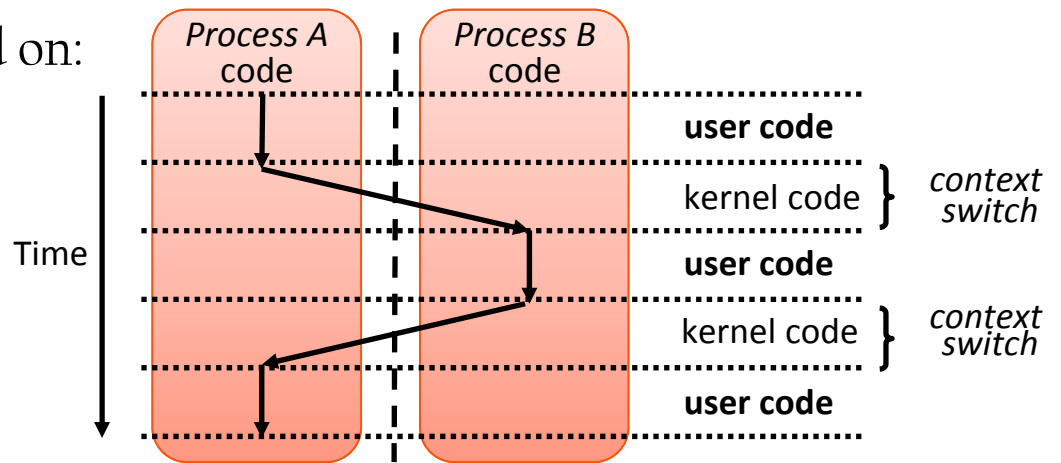
- piece of OS code handling access to CPU

- goes from one process to another
 - to use CPU time most efficiently...
 - and be fair to all processes (time sharing)

- via a *context switch*

- decision based on:

- » fairness
- » priority
- » events
(mouse, key,
timer, ...:
interrupts)



Example of Concurrency Issues

Assume shared memory

- processes can read/write the same variable

Start $f()$ and $g()$ at the “same” time (and $x=0$)


$f()$:

```
global x
for i=1 to 1000
  x = x + 1
```

$g()$:

```
global x
for i=1 to 1000
  x = x + 1
```

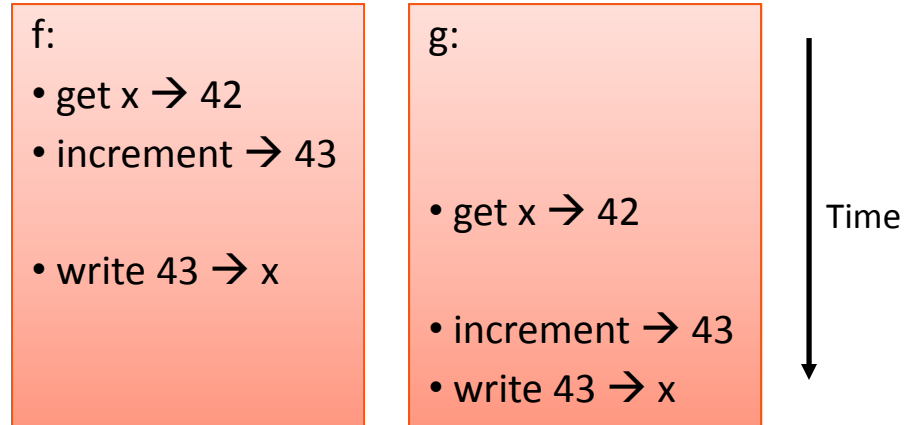
- “ $x=x+1$ ” means read memory, add 1, write to memory
- remember, execution will lead to intertwined code lines
- question: what is x when all is finished?



Execution Order

Suppose $x = 42$

- at a certain point



Net result: lost an increment

- called a **race condition** (non-deterministic results)
- hard to detect: problems are intermittent
 - all possible schedules have to be safe
 - but number of possible schedule permutations is huge!

Safe way to increment x without issues?



Synchronization

Concurrently-executing threads/processes should not execute **specific portions of a program** at the same time

- notion of critical sections (CS)
 - $x=x+1$ should be done atomically
- could make special hardware instructions
 - but does not extend to database systems or web servers
- need software solutions
 - used by the coder, or/and by the kernel



Critical Section Requirements

Mutual Exclusion

- only one process is in CS at any time

Progress

- not every process waiting for CS indefinitely
 - no deadlock

Bounded Waiting

- bounded #times others get CS while one is waiting
 - no livelock

When two trains approach each other at a crossing, both shall come to a full stop and neither shall start up again until the other has gone.

Statute passed by Kansas Legislature

Also: fairness, simplicity/efficiency, ...



Taking Turn?

Use a variable
to determine turn

```
Process P0:  
repeat  
  while(turn!=0){};  
  CS  
  turn:=1;  
  <whatever>  
forever
```

```
Process P1:  
repeat  
  while(turn!=1){};  
  CS  
  turn:=0;  
  <whatever>  
forever
```

- achieves Mutual Exclusion (thru busy wait)
- **but** progress requirement is **not** satisfied
 - it requires strict alternation of CS's.
 - If a process asks for CS more often than the other, it can't get it



Checking Intent?

Initialization:
flag[0]:=flag[1]:=false

Use 2 intent flags

```
Process P0:
repeat
  flag[0]:=true;
  while(flag[1]){};
  CS
  flag[0]:=false;
  <whatever>
forever
```

```
Process P1:
repeat
  flag[1]:=true;
  while(flag[0]){};
  CS
  flag[1]:=false;
  <whatever>
forever
```

- Mutual Exclusion is satisfied
- **but** not the progress requirement
 - for the (interleaved) sequence:
 - » flag[0]:=true
 - » flag[1]:=true
 - both processes will wait forever (deadlock)



Peterson's Solution

Tie-breaking

- Even if both flags go up
 - and no matter how the instructions are interleaved
- ... turn will always end up as either 0 or 1

```
Initialization:  
flag[0]:=flag[1]:=false  
turn:= 0 or 1
```

```
Process P0:  
repeat  
    flag[0]:=true;  
    // 0 wants in  
    turn:= 1;  
    // 0 gives a chance to 1  
    while (flag[1]&turn==1);  
        CS  
    flag[0]:=false;  
    // 0 is done  
    <whatever>  
forever
```

```
Process P1:  
repeat  
    flag[1]:=true;  
    // 1 wants in  
    turn:=0;  
    // 1 gives a chance to 0  
    while (flag[0]&turn==0);  
        CS  
    flag[1]:=false;  
    // 1 is done  
    <whatever>  
forever
```



Correctness

Mutual exclusion holds since:

- both P_0 and P_1 in CS means
 - both `flag[0]` and `flag[1]` must be true...
 - and `turn=0` and `turn=1` (at same time): impossible

```
Process P0:
repeat
  flag[0]:=true;
  // 0 wants in
  turn:= 1;
  // 0 gives a chance to 1
  while (flag[1]&turn=1);
    CS
  flag[0]:=false;
  // 0 is done
  RS
forever
```

Progress

- P_0 can be kept out of CS only if stuck in while loop
 - `flag[1] = true` and `turn = 1`.
- If P_1 not ready to enter CS, `flag[1] = false` → P_0 can enter CS
- If P_1 has set `flag[1]`, it is also in its while loop;
then either P_0 or P_1 will go depending on value of `turn`
- progress condition is met



Correctness

Bounded Wait?

- suppose P_0 gets to go this time;

Can it go a second time indept of what P_1 does?

- if P_1 enters CS , then $turn=1$
 - but will then reset $flag[1]$ to false on exit
 - » thus allowing P_0 to enter CS
- what if P_1 tries again, and has time to reset $flag[1]=true$ before P_0 gets to its CS?
 - it must also set $turn=0$
 - » since P_0 is (stuck) past the point where it sets $turn$ to 1:
 - P_0 will get to enter CS after at most one CS entry by P_1

```
Process P0:
repeat
  flag[0]:=true;
  // 0 wants in
  turn:= 1;
  // 0 gives a chance to 1
  while (flag[1]&turn=1);
    CS
  flag[0]:=false;
  // 0 is done
  RS
forever
```



Semaphores

Introduced by Dijkstra (1968)

- higher-level tool for process synchronization
 - uses an integer variable s ...
 - initialized with # of available resources (think: room reservation; license counts)
 - and two operations:
 - $P(s)$:
atomically executes $s := s-1$
delays until $s > 0$

until there's an opening
 - $V(s)$:
atomically executes $s := s+1$
- binary semaphore = mutex
 - $P(s)$ before CS, $V(s)$ after CS



No Busy Waiting!

queue is FIFO
or priority based

OS semaphore w/ wait queue

- to put process to sleep

```
typedef struct semaphore {  
    int value;  
    SleepingProcessesList L;  
} Semaphore;
```

```
void P(Semaphore *S) {  
    S->value = S->value - 1;  
    if (S->value < 0) {  
        add this process to S.L;  
        block();  
    }  
}
```

```
void V(Semaphore *S) {  
    S->value = S->value + 1;  
    if (S->value >= 0) {  
        remove process P from S.L;  
        wakeup(P)  
    }  
}
```

- block(): remove process from active pool
- wakeup(): back in kernel pool and resume



Possible Use of Concurrency

Many recursive methods involve 2 calls

- e.g., merge sort: left half and right half
- can be co-invoked!


Creation of threads by, e.g., fork/join

- `fork()` spawns copy





Forking

BEFORE



AFTER

Parent
pid = 521



Child
pid = 0

Next line of code is typically: `if (pid==0) [...] else [...]`



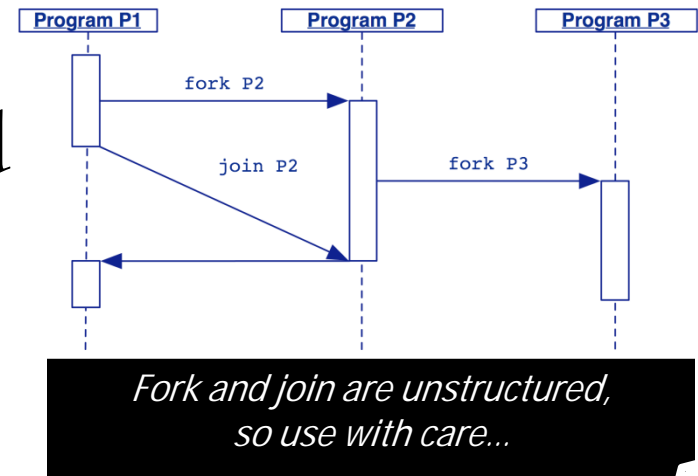
Possible Use of Concurrency

Many recursive methods involve 2 calls

- e.g., merge sort: left half and right half
- can be co-invoked!

Creation of threads by, e.g., fork/join


- `fork()` spawns copy
- `join()` waits for it to end



Communication & Synchronization


We looked at the shared memory model

- proc. “communicate” thru shared variables
- requires synchronization for memory access



Other approach: message-passing



- *communication & synchronization* are combined
- (a)synchronous comm.




Why Should You Care?

Parallel computing increasingly important

- was trendy in the 70's and 80's
- but *really* necessary now... why?
- Moore's law, with a twist



my old, old laptop, circa 2010




my old laptop, circa 2012



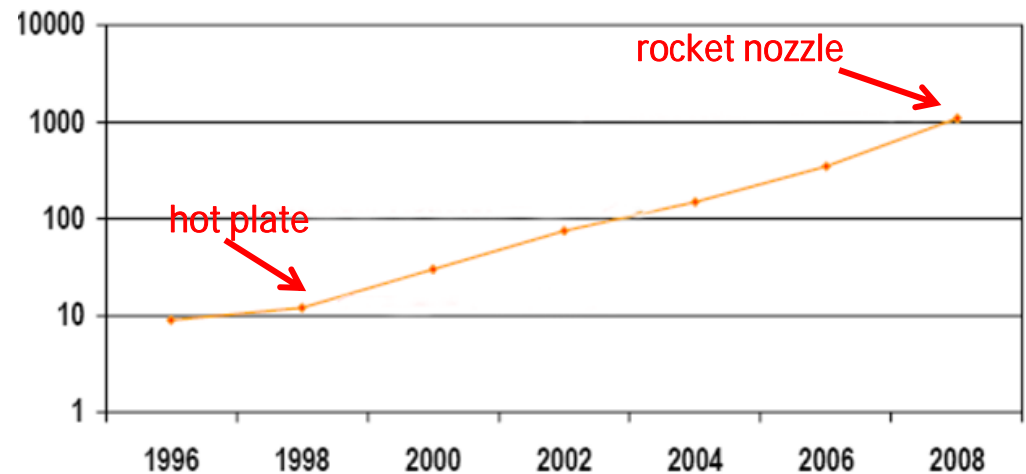
Reason: Too Hot...

Increasing transistor density...
... plus increasing clock-speed ...
means increasing *power density*

■ parallelism
to the rescue



Power Density
(Watts/cm²)



source: Shekhar Borkar, Intel

