

CS 2

Introduction to Programming Methods



Last Time


Concurrency and Parallelism

- synchronization of multiple threads

Reason: Too Hot...

Increasing transistor density...
... plus increasing clock-speed ...
means increasing *power density*

- parallelism to the rescue



source: Shodor, Fortis, Intel

CS2 - INTRODUCTION TO PROGRAMMING METHODS

21



CS2 - INTRODUCTION TO PROGRAMMING METHODS

2

Today's Lecture

Numerics: when good math goes wrong

Computer programmers are fallible

- losing satellite due to public vs. private var.
 - self.setValue(.) bypassed; worked at the time...
 - but code changed, so value overwrote other data
- losing a Mars orbiter due to conversion
 - we have switched to the metric system, right?

Computers are perfect with numbers...

- right?



Integer Paradise

Dealing with bounded ints is great

- exact computations easy
 - as long as integers below a prescribed value (often, 32 bits)
 - 2 bits: 0 to 3
 - 8 bits: 0 to 255
 - n bits: 0 to $2^n - 1$
 - signed int
 - most commonly: two's complement
 - why?

$$-x \equiv \bar{x} + 1$$

$$v = -d_{31}2^{31} + \sum_{n=0}^{30} d_n 2^n \quad \gg \text{add/sub painless (bit overflow ignored)}$$

Most significant bit

0	1	1	1	1	1	1	1	=	127
0	1	1	1	1	1	1	0	=	126
0	0	0	0	0	0	0	1	=	2
0	0	0	0	0	0	0	0	=	1
0	0	0	0	0	0	0	0	=	0
1	1	1	1	1	1	1	1	=	-1
1	1	1	1	1	1	1	0	=	-2
1	0	0	0	0	0	0	1	=	-127
1	0	0	0	0	0	0	0	=	-128



What About Reals?

Not so simple...

```
Python 2.5.1
Type "help", "copyright", "credits" or "license" for more information.
>>> 0.1
0.10000000000000001
>>>
```

- what's wrong?? [by the way, newer versions say: 0.1]

Tenths not very easy to represent in binary

- $0.1 = 0b0.000110011001100110011001100110011001100110011001100110011010\dots$

So, computations often not perfect

- and we can send people to the moon??

```
>>> sum = 0.0
>>> for i in range(10):
...     sum += 0.1
...
>>> sum
0.9999999999999999
```



Fixed-Point Representation

Approximate positive reals w/ powers of 2

- first, decide how many bits to use
 - 32 bits or 64 bits quite common
- then, decide where to put the “binary point”
 - for instance, point before last bit → quantum of 0.5

Signed reals?

- same two's complement idea
 - e.g., 32 bits for numbers between -1 and $1 - 2^{-31}$

$$v = -d_0 2^0 + \sum_{n=1}^{32} d_n 2^{-n}$$

$$b_i \ b_{i-1} \ \dots \ b_2 \ b_1 \ b_0 . b_{-1} \ b_{-2} \ b_{-3} \ \dots \ b_{-j} \quad \sum_{k=-j}^i b_k \cdot 2^k$$



Adaptive Representation?

Fixed-point is fast and simple

- but limited!
 - fixed accuracy, quite limited range
 - in fact, trade precision for range

Large range and high precision?

- would require lots of bits
- potentially wasteful



Floating-Point Representation

IEEE 754 definition



Variants:

- single precision: 8 exp bits, 23 mant . bits
 - 32 bits total
- double precision: 11 exp bits, 52 mant . bits
 - 64 bits total
- extended precision: 15 exp bits, 63 mant . bits
 - mostly in Intel-compatible machines; stored in 80 bits
 - > 1 bit wasted



Single Precision FP


For 32 bits



■ represents a value of $(-1)^s (1.m)_2 2^{e-127}$

■ $1 \leq (1.m)_2 < 2$ $0 < e < 255$

■ example: 0xD1500000



■ value = $-1.625 \cdot 2^{35} = -5.5834e+10$

note: hexadecimal notation

xxxx: 0, 1, 2, ..., 8, 9, A, B, C, D, E, F



Special Numbers

s	e	m	value
0	all zeros	all zeros	0
1	all zeros	all zeros	-0
0	all ones	all zeros	∞
1	all ones	all zeros	$-\infty$
0 or 1	all ones	non-zero	NaN
0 or 1	all zeros	non-zero	$(-1)^s (0.m)_2 2^{-126}$

Check: max = $3.402823466e+38$

min = $1.401298464e-45$

Examples: $0 \times \infty$,
 $\text{sqrt}(-1)$, ...

gradual underflow



Double-Precision FP

Same exact principles

- represents value of $(-1)^s (1.m)_2 2^{e-1023}$

max = 1.7976931348623157e+308

min = 5e-324



Floating-Point Operations

Conceptually

- First, compute exact result
- Make it fit into desired precision
 - Possibly overflow if exponent too large
 - Possibly round to fit into mantissa

Example: $(-1)^{s_1} m_1 2^{e_1} \times (-1)^{s_2} m_2 2^{e_2}$

- exact result: $(-1)^s m 2^e$
 - $s = s_1 + s_2$, $m = m_1.m_2$, $e = e_1 + e_2$
 - $m \geq 2$? shift m right, increment e ; e out of range? overflow
 - round m to fit mantissa precision



FP Addition


Operands: $(-1)^{s_1} m_1 2^{e_1}$ and $(-1)^{s_2} m_2 2^{e_2}$

- Assume $e_1 > e_2$

exact result: $(-1)^s m 2^e$

- exponent $e=e_1$; sign s ; significand m

- result of signed align & add



Normalization:

- If $m \geq 2$, shift m right, increment e
- if $m < 1$, shift m left k positions, decrement e by k
- overflow if e out of range, round m to fit precision



Math vs. Numerics

Properties of Addition

- commutative? YES
- associative? NO
 - overflow and rounding
- 0 is additive identity? YES
- always additive inverse ALMOST
 - except for $\pm\text{infinity}$ & NaNs
- $a \geq b \Rightarrow a+c \geq b+c$? ALMOST
 - except for $\pm\text{infinity}$ & NaNs



[Numerics in C]

C provides two levels

- float single precision
- double double precision

Casting between int, float, & double

- double or float to (64 bit) int
 - truncates fractional part (rounding toward zero)
- int to double
 - exact conversion
- int to float
 - depends on rounding mode...




More Numerical Catastrophes...

Understanding FP helps with accuracy...

- $p(\cdot)$ polynomial having $x=1, 2, \dots, 25$ as roots
- unique one of degree 25
- two examples of eval:

$$p(x) = \prod_{i=1}^{25} (x - i)$$



$p(x) = -1551121004330985984000000 +$
 $59190128811701203599360000 x -$
 $10048017154835116154880000 x^2 +$
 $102339530601744675672576000 x^3 -$
 $70874145319837672677196800 x^4 +$
 $35770355645907606826362624 x^5 -$
 $13746468217967926978680000 x^6 +$
 $4144457803247115877036800 x^7 -$
 $1001369304512841374110000 x^8 +$
 $196928100451110820242880 x^9 -$
 $31882014375298512782500 x^{10} + 4284218746244111474800$
 $x^{11} - 480544558742733545125 x^{12} + 45145946926994481865$
 $x^{13} - 3557372853474553750 x^{14} + 234961569422786050 x^{15} -$
 $12972753318542875 x^{16} + 595667304367135 x^{17} -$
 $22563937825000 x^{18} + 696829576300 x^{19} - 17247104875 x^{20} +$
 $333685495 x^{21} - 4858750 x^{22} + 50050 x^{23} - 325 x^{24} + x^{25}$



Numerics Can Ruin a Game Too

Depth ac



How Numerics Can Kill



1996: Ariane 5 explodes

- problem quickly identified as numeric
 - but this part worked just fine on Ariane 4!
- Error was due to “a data conversion from *64-bit floating point* to *16-bit signed integer* value. The floating point number which was converted had a value greater than what could be represented by a 16-bit signed integer.”
 - Ariane 5 was... faster
- \$7.5B in development and launch lost

