

CS 2

Introduction to Programming Methods



Last Time

Numerics & FFT


- Fourier transform in $O(n \log n)$



Signal Processing: Example

Edit image by altering frequency content

- can also do:
 - Darth Vader voice
 - (de)noising
 - (de)blurring
 - compression
 - etc...



CS2 - INTRODUCTION TO PROGRAMMING METHODS

19



CS2 - INTRODUCTION TO PROGRAMMING METHODS

Artificial Intelligence

Overused, overloaded, & misused term

- intelligence is the computational part of the ability to achieve goals
- AI: science and engineering of making intelligent machines
 - especially, intelligent computer programs



Artificial Intelligence

Quick history

- coined in 1956 by McCarthy (teacher!) & co
 - ideas dating back to Alan Turing in the 50's
 - logic programming used towards electronic brain
- a bit too optimistic initially...
- but interesting developments since then
 - '97: deep blue beats Kasparov, a chess grandmaster
 - '05-'07: DARPA Grand/Urban challenge
 - 100+/50+ miles of autonomous driving in desert/city
 - '11: Watson beats masters of Jeopardy!



Wide Range of AI Research

Typical subfields of AI

- problem solving & deduction/reasoning
 - puzzles, logic
- planning & coordination
 - autonomous agents to achieve global goal(s)
- machine learning
 - infer relationships, find patterns from data
- natural language processing
 - reading and understanding human language(s)
- knowledge representation
 - how do you encode the obvious?
 - » “the flesh is weak” does not mean that the meat is on sale



AI for Games I

For basic puzzles, easy

- search graph of possibilities

- BFS or DFS

- small size graphs for:

- 8-puzzle, sudoku, ...

- often impossible: too big...



- needs heuristic—e.g., prune off branches of the search

- example: A* algorithm

- “best”-first search

- » “prioritized” Dijkstra’s

- » mix btw BFS & DFS



AI for Games II

2-Player Games?


- Min-max (minimax) algorithm predominant
 - always assume the adversary will play best move
 - search for best move *within this assumption* (and for n moves ahead)
 - dual recursion [*minimizing the max loss; or maximize the min gain*]
 - again, possibly with heuristics (e.g., AlphaBeta)
 - ex: tic-tac-toe

Computer's turn

- maximize outcome

Adversary's turn

- minimize outcome



AI for Games II

2-Player

■ Min

■ al

■ se



■ ex

Computer's turn

- maximize

Adversary's turn

- minimize

```
function minimax(node, depth, maximizingPlayer)
  if depth = 0 or node is a terminal node
    return the heuristic value of node

  if maximizingPlayer
    bestValue := -∞
    for each child of node
      v := minimax(child, depth-1, false)
      bestValue := max(bestValue, v)
    return bestValue
  else //minimizing player
    bestValue := +∞
    for each child of node
      v := minimax(child, depth-1, true)
      bestValue := min(bestValue, v)
    return bestValue
```

nant

ove

n moves ahead)

min gain]

can play

What opponent
could reply



AI for Games II

2-Player Games?

OR

```


function minimax(node, depth, ply)
  if node is terminal or if depth <= 0:
    return heuristic evaluation of node * ply
   $\alpha = -\infty$ 
  for child in node
    // max(a,b)=-min(-a,-b)
     $\alpha = \max(\alpha, -\text{minimax}(\text{child}, \text{depth}-1, -\text{ply}))$ 
    // find max if I am the player, min otherwise
  return  $\alpha$ 
  
```

Computer's turn

- maximize outcome

Adversary's turn

- minimize outcome



AI for Games III

State of the art

- easy: 8-puzzle, Tic-tac-toe, Connect Four.
- thought hard, now easy: Checkers, Chess, Backgammon.
 - “easy” meaning easily beating humans
- still hard: Go, Othello.



AI for Behavior

Finite state machine

- abstract model of behavior

- starting state
- event-based directed graph

- autonomous agents

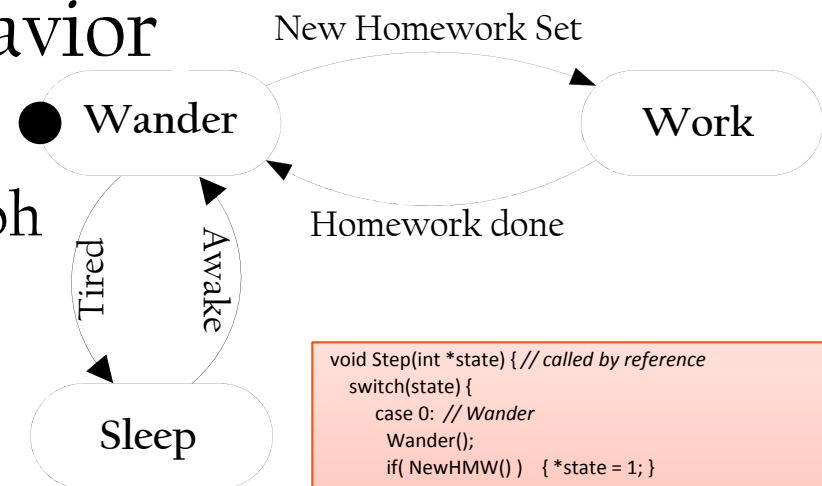
- e.g., in games

- automata theory

- various types
 - (non)deterministic, pushdown, ...
- can recognize languages too!

You also need to:

- ✓ eat & have fun
- ✓ attend CS2



```
void Step(int *state) { // called by reference
    switch(state) {
        case 0: // Wander
            Wander();
            if( NewHmw() ) { *state = 1; }
            if( Tired() ) { *state = 2; }
            break;
        case 1: // Work
            Work();
            if( HmwDone() ) { *state = 0; }
            break;
        case 2: // Sleep
            Sleep();
            if( Awake() ) { *state = 0; }
            break;
    }
}
```



Neural Networks



Motivation: analogy to the brain

- massively interconnected neurons

Artificial neuron?

- much simpler, abstract notion
 - “processing unit”, receiving and sending information




- connect them to do complex computations
 - connected through a weighted graph
 - parallel processing



Simple Neuron (Perceptron)

Two inputs, one output



- $y = f(w_1x_1 + w_2x_2)$
- activation function f is, e.g., step function:
 - $f(x) = 1$ for $x \geq \theta$, 0 otherwise




Examples of Simple Perceptrons

Binary gates

- {in/out}put 0 or 1


OR

X_1	X_2	Y
0	0	0
1	0	1
0	1	1
1	1	1




AND

X_1	X_2	Y
0	0	0
1	0	0
0	1	0
1	1	1



AND NOT



X_1	X_2	Y
0	0	0
1	0	1
0	1	0
1	1	0



What About XOR?

Apparent glitch...


- no weights will do
- why? not linearly separable




XOR

x_1	x_2	Y
0	0	0
1	0	1
0	1	1
1	1	0

- but by adding a “hidden layer”, easy




Simply implementing:
$$x_1 \text{ XOR } x_2 = (x_1 \text{ AND NOT } x_2) \text{ OR } (x_2 \text{ AND NOT } x_1)$$



Training Multilayer Networks

Perceptrons can be trained

- tweak weights and thresholds
 - or more complex parameterized activation functions
- to “classify” inputs properly
 - feed inputs and look at outputs
 - not right? change weights and thresholds
 - and try again
- use backpropagation, reinforcement learning, evolutionary algorithms...
 - recognizing cats in videos w/ 9 layers, 1B connections...



More Neural Networks

Talked only about feed-forward



- unidirectional flow

Recurrency in networks too

- “feedback” loop
- Hopfield, Elman, ...
- can now handle sequences
 - “states”, or memory

See also Bayesian networks

- probabilistic graphical model
-



FYI



John Hopfield
at Caltech '80-'97

