



Assignment 6 Recitation

Concurrency



The Threading Model

- Multiple lines of execution, one set of shared memory
- As opposed to a process
- To handle this without errors, we require...



Concurrency Primitives

- Focus on three types:
 - Threads
 - Mutexes
 - Semaphores
- Already covered high level in lecture
- Now built into C++11
- Makes your life easier! (as compared to POSIX or other libraries)



Threads

- With C++11, threads can run any function
- The function name is the first input to the thread constructor
- The function arguments follow, in the same order in the constructor as in the function



Threads

- An example:

```
void thread_adder(int a, int b) {  
    std::cout << a << std::endl;  
    std::cout << b << std::endl;  
    std::cout << a + b << std::endl;  
}
```

```
int a = 1;  
int b = 2;  
std::thread t(thread_adder, a, b);  
t.join()
```



Threads

- A side note: joining
 - Blocks execution on the current thread until the joined thread has finished
 - As opposed to detaching, where both threads are free to run simultaneously



Mutexes

- Use when only one thread should use a resource at a time

```
std::mutex* m = new std::mutex();    // Create a mutex.
```

```
m->lock();                          // Lock mutex.
```

```
m->unlock();                         // Unlock mutex.
```



Mutex Example

```
void a(std::mutex* m) {  
    m->lock();  
    std::cout << "In a!" << std::endl;  
}
```

```
//Lock shared by a, b  
std::mutex m;  
a(&m);  
b(&m);
```

```
void b(std::mutex* m) {  
    m->lock();  
    std::cout << "In b!" << std::endl;  
    m->unlock();  
}
```




Deadlock!

```
void a(std::mutex* m) {  
    m->lock();  
    std::cout << "In a!" << std::endl;  
}
```

```
//Lock shared by a, b  
std::mutex m;  
a(&m);  
b(&m);
```

```
void b(std::mutex* m) {  
    m->lock();  
    std::cout << "In b!" << std::endl;  
    m->unlock();  
}
```

Deadlock fixed!

```
void a(std::mutex* m) {  
    m->lock();  
    std::cout << "In a!" << std::endl;  
    m->unlock();  
}  
  
void b(std::mutex* m) {  
    m->lock();  
    std::cout << "In b!" << std::endl;  
    m->unlock();  
}
```

```
//Lock shared by a, b  
std::mutex m;  
a(&m);  
b(&m);
```



Another Mutex Example

```
void fill(std::vector<int>* vec) {  
    for (int i = 0; i < 2000; ++i) {  
        vec->push_back(i);  
    }  
}
```



Another Mutex Example

```
vector<int> nums;  
std::thread a(fill, &nums);  
std::thread b(fill, &nums);  
a.join();  
b.join();  
std::cout << "size of nums: " << nums.size() << std::endl;
```



Race conditions!

```
vector<int> nums;  
std::thread a(fill, &nums);  
std::thread b(fill, &nums);  
a.join();  
b.join();  
std::cout << "size of nums: " << nums.size() << std::endl;
```



A quick fix

```
std::mutex lock;
```

```
...
```

```
void fill(std::vector<int>* vec, std::mutex* m) {  
    for (int i = 0; i < 2000; ++i) {  
        m->lock();  
        vec->push_back(i);  
        m->unlock();  
    }  
}
```



Semaphores

- Can be thought of as generalized mutexes
 - Instead of only 0 or 1 threads holding the lock, 0..n threads can hold the lock
- Not available in C++11 :(

```
Semaphore* s = new Semaphore(2);
```

```
s->dec();
```

```
s->dec(); // Semaphore count is now 0, more decs will block
```

```
s->inc(); // Semaphore now has space
```

```
s->dec();
```

```
s->value(); // Returns 0
```



Consumer Producer Queues

- Some threads produce data
- Some threads read/consume data
- These consumers and producers need to share a fixed size buffer
- Two potential problems
 - Empty buffer, consumer must wait
 - Full buffer, producers must wait
- Apply semaphores?




Consumer Producer Queues

- What do we need to track?
 - Number of slots filled.
 - Number of slots empty
 - Use a semaphore for each value
- Producers
 - Decrement empty slots semaphore
 - Add value
 - Increment filled slots semaphore
- Consumers
 - Decrement filled slots semaphore
 - Use value
 - Increment empty slots semaphore

Dining philosophers

- 5 silent hungry philosophers
- One bowl of spaghetti
- 5 forks placed as shown
- Philosophers alternate eating and thinking
- Must have both forks to left+right to eat
- Cannot grab forks for each other
- Replace forks after eating
- How to design behavior such that each philosopher won't starve?





Proposal

- Think until the right fork is free; when it is, pick it up
- Think until the left fork is free; when it is, pick it up
- When both forks are held, eat for a fixed amount of time
- Then, put the left fork down
- Then, put the right fork down
- Repeat



Deadlock

- All start picking up the right fork at the same time
- Wait forever for the left fork to be ready!



Round Robin


- Since they are silent, imagine a third party helper
- Tells the philosophers when to eat
- Only allows one philosopher to eat
- Deadlock?
- Implement using primitives mentioned above




Homework notes:

- Each philosopher runs in its own thread!
- Every philosopher must eventually eat!


Raytracing!



Raytracing!



Raytracing!





Questions?